

SAKSHI S. MORE

Senior Software Engineer | Full Stack Developer | AI/GenAI Engineer | Python | React | Node.js | AWS

Austin, TX | +1 210-473-7294 | sakshismore.dev@gmail.com | [linkedin.com/in/sakshisandeep/](https://www.linkedin.com/in/sakshisandeep/) | github.com/Sakshimistic

PROFESSIONAL SUMMARY

Results-driven **Software Engineer** with 5+ years of experience architecting and delivering **scalable full-stack applications**, **distributed backend systems**, and **cloud-native microservices** across enterprise and research environments. Deep expertise in **Python, TypeScript, and JavaScript** with production delivery across **React, Next.js, Node.js, Django, and FastAPI**. Hands-on experience integrating **LLM-powered features**, **RAG pipelines**, and **AI/ML model serving** using OpenAI API, LangChain, and vector search systems. Proven track record optimizing system performance at scale, leading CI/CD automation, and shipping production-grade features in **Agile** environments.

SKILLS

Languages: Python, JavaScript, TypeScript, Java, SQL, Bash/Shell, R, Scala

Frontend: React, Next.js, Redux, Tailwind CSS, HTML5, CSS3, Webpack, Vite, REST APIs, GraphQL

Backend & Frameworks: Node.js, Express.js, FastAPI, Django, Flask, Microservices, WebSockets, gRPC, Apache Kafka

AI / ML & GenAI: LangChain, OpenAI API, RAG (Retrieval-Augmented Generation), Prompt Engineering, Fine-Tuning, Embeddings, Vector Search, Agent Orchestration, Multi-Agent Systems, TensorFlow, PyTorch, Scikit-learn, Hugging Face Transformers, XGBoost, FAISS, Pinecone, Weaviate

Databases & Storage: PostgreSQL, MySQL, MongoDB, DynamoDB, Redis, Snowflake, Elasticsearch, FAISS, Azure Cosmos DB, SQLAlchemy

Cloud & DevOps: AWS (EC2, S3, Lambda, SageMaker, RDS, CloudWatch), GCP, Microsoft Azure (Azure ML, Azure OpenAI, AKS, Cosmos DB), Terraform, Docker, Kubernetes, GitHub Actions, Jenkins, CI/CD Pipelines, Linux/Unix

MLOps & Observability: MLflow, Kubeflow, AWS SageMaker, Vertex AI, Model Monitoring, Drift Detection, Prometheus, Grafana, Datadog, Streamlit, Plotly

Certifications: Oracle Data Science Professional | Oracle Generative AI Professional | Oracle AI Foundations Associate

WORK EXPERIENCE

Software Research Engineer

Feb 2023 – Present

University of Texas, Austin, TX

- Designed and owned a **distributed semantic search system** in Python and FastAPI processing 3M+ research records, leveraging **FAISS vector indexing** and **LLM-generated embeddings** (OpenAI API) to reduce query latency from 4.2s to under 800ms — a 5x performance improvement.
- Built an end-to-end **RAG pipeline** using LangChain and OpenAI GPT-4 enabling researchers to extract key findings from 500-page document corpora in under 8 seconds; adopted across 3 federally funded grant projects.
- Architected **serverless data ingestion workflows** on **AWS Lambda and S3** with event-driven triggers, cutting infrastructure overhead by 55% and scaling to 10K+ concurrent sessions without throttling.
- Engineered a **multi-stage CI/CD pipeline** using **GitHub Actions, Docker**, and containerized test environments achieving a 99.6% build success rate across 12 repositories; reduced release cycle time by 70%.
- Built real-time experiment monitoring dashboards using **Plotly and Streamlit** connected to live **PostgreSQL** streams, replacing 5 manual reporting workflows and recovering 18+ person-hours per sprint.
- Integrated **vector similarity search (FAISS + Weaviate)** to power a document recommendation engine with sub-second retrieval, increasing researcher engagement by 40% versus keyword-only search.

Software Engineer Intern

May 2023 – Aug 2023

Granite Construction, CA, USA

- Engineered 14 **RESTful API endpoints** in **Node.js and Express.js** backed by **PostgreSQL**, enabling real-time field project status logging; reduced data turnaround time by 3x over prior spreadsheet workflows.
- Developed a **React/Redux** asset-tracking module providing live visibility into \$50M+ of active construction equipment across 6 regional sites, cutting asset-location resolution from hours to under 2 minutes.
- Automated end-to-end deployment pipelines with **Docker and GitHub Actions**, compressing a 3-day manual release to a 2-hour automated flow and eliminating 4 recurring production incidents.
- Optimized high-traffic **MySQL** reporting queries via execution plan analysis and composite index tuning, reducing finance dashboard load time by 47%; enabled weekly cost-variance reporting for a 30-person PM org.
- Delivered 4 production-grade features within a 14-week internship cycle; contributed to a 22% velocity improvement measured across the final 3 Agile sprints.

Full-Stack Software Developer

Jun 2021 – Nov 2022

Tata Consultancy Services, Pune, India

- Led full-stack development of a **policyholder self-service portal** using **React, Next.js, and Django REST Framework**, serving 200K+ active users across 3 markets; reduced inbound support ticket volume by 34% post-launch.
- Re-architected a monolithic data layer by migrating critical datasets to **MongoDB with Azure Cosmos DB** for multi-region replication, improving read throughput by 60% and sustaining a 99.9% uptime SLA under peak load.
- Containerized and orchestrated 5 production microservices on **Azure Kubernetes Service (AKS)**, enabling zero-downtime quarterly deployments; reduced incident rollback time from 4 hours to under 15 minutes.
- Centralized observability with **Elasticsearch-backed structured logging**, proactively identifying 3 P1 performance bottlenecks that had degraded API throughput by 30% — all resolved before SLA breach.
- Partnered with business analysts, QA, and client stakeholders to ship a 17-feature roadmap on schedule, contributing to a 22-point increase in client satisfaction score over two consecutive quarters.

Software Engineer

Jan 2021 – May 2021

Mindtree, Pune, India

- Automated financial data ETL across 12 source systems using **Python, Pandas, and Bash**, eliminating 40 hours of manual monthly reconciliation and reducing processing errors by 38%.
- Built a **Flask REST API** backed by MySQL to serve cleansed records to internal reporting tools, enabling a 25-person ops team to self-serve data without analytics team routing.
- Consolidated 6 fragmented executive reporting workflows into a single automated pipeline using **Matplotlib and Seaborn**, cutting weekly KPI report prep from 6 hours to under 45 minutes.

EDUCATION

Master of Science, Computer Science (Data Science) | University of Texas at San Antonio

May 2025

Bachelor of Engineering, Information Technology | Dr. D.Y. Patil College of Engineering, India

May 2021

PROJECTS

Job Hunt Agent | TypeScript, Node.js, AI Agent Orchestration, REST APIs

- Built an autonomous **AI-powered job search agent** using TypeScript that scrapes job boards, matches listings against a candidate profile using LLM-based semantic scoring, and drafts tailored cover letters — reducing manual job-hunting effort by 80%.
- Implemented multi-step **agent orchestration** with tool-calling loops, retry logic, and structured JSON outputs using the OpenAI API; deployed on Replit for always-on operation.

Log Classification System | Python, NLP, LLM, Regex, Jupyter

- Designed a **hybrid log classification framework** combining Regex pattern matching, traditional NLP (TF-IDF + classification), and LLM-based zero-shot inference to handle varying complexity in production log patterns.
- Achieved 94%+ classification accuracy on real-world log datasets; architecture is model-agnostic and supports **plug-and-play LLM backends** (OpenAI, local Ollama models).

Twitter Real-Time Streaming Pipeline | Python, Apache Kafka, Elasticsearch, Docker, Kibana

- Built a **real-time data streaming pipeline** ingesting live Twitter data via the Twitter API, processing events with **Apache Kafka**, storing in **Elasticsearch**, and visualizing engagement trends in Kibana — fully containerized with Docker Compose.
- Designed producer/consumer architecture supporting 1,000+ events/sec with schema validation and dead-letter queuing for fault tolerance.

SimpleSocialAPI | FastAPI, SQLAlchemy (Async), PostgreSQL, JWT, Streamlit, ImageKit

- Engineered a full-stack **social media REST API** with **JWT-based authentication**, async SQLAlchemy ORM, media uploads via ImageKit CDN, and a Streamlit frontend; implemented pagination, rate limiting, and OpenAPI docs.
- Applied **async/await patterns** throughout the stack, reducing average API response time by 35% compared to synchronous baseline.

Retail Customer Behavior Analytics | Python, Pandas, Scikit-learn, Plotly, Jupyter

- Performed end-to-end **customer segmentation and churn prediction** on retail transaction data using RFM analysis and XGBoost; produced interactive Plotly dashboards surfacing actionable business insights.
- Applied SHAP explainability to the churn model, increasing stakeholder trust and enabling targeted retention campaigns with projected 15% reduction in customer churn.

Flask Travel API | Python, Flask, REST API, PostgreSQL

- Built a **RESTful travel itinerary API** with CRUD operations, user authentication, and location-based search; documented with Swagger UI and deployed with Docker.